

CSE3013- ARTIFICIAL INTELLIGENCE

TOPIC- CREDIT CARD FRAUD DETECTION

Team Members-

Ananya Pantvaidya (20BCE0678)

Arsh Ansari (20BCE0371)

Kushagra Srivastava (20BCE2060)

Priadarshni Muthukumar (20BCE2510)

Importing the necessary libraries

```
In [1]: import numpy as np
import pandas as pd
import seaborn as sns
import matplotlib.pyplot as plt

import warnings
warnings.filterwarnings('ignore')
```

Importing the dataset

We have downloaded our dataset from Kaggle

```
In [2]: data=pd.read_csv("creditcard.csv")
```

Taking a look at the features of the dataset

```
In [3]: print(data.columns)

Index(['Time', 'V1', 'V2', 'V3', 'V4', 'V5', 'V6', 'V7', 'V8', 'V9', 'V10',
       'V11', 'V12', 'V13', 'V14', 'V15', 'V16', 'V17', 'V18', 'V19', 'V20',
       'V21', 'V22', 'V23', 'V24', 'V25', 'V26', 'V27', 'V28', 'Amount',
       'Class'],
      dtype='object')
```

V1-V28 are the result of a PCA dimensionality reduction


```
In [4]: data.head()
```

```
Out[4]:
```

	Time	V1	V2	V3	V4	V5	V6	V7	V8	
0	0.0	-1.359807	-0.072781	2.536347	1.378155	-0.338321	0.462388	0.239599	0.098698	0.363

	Time	V1	V2	V3	V4	V5	V6	V7	V8	
1	0.0	1.191857	0.266151	0.166480	0.448154	0.060018	-0.082361	-0.078803	0.085102	-0.255
2	1.0	-1.358354	-1.340163	1.773209	0.379780	-0.503198	1.800499	0.791461	0.247676	-1.514
3	1.0	-0.966272	-0.185226	1.792993	-0.863291	-0.010309	1.247203	0.237609	0.377436	-1.387
4	2.0	-1.158233	0.877737	1.548718	0.403034	-0.407193	0.095921	0.592941	-0.270533	0.817

5 rows × 31 columns




In [5]:

```
data.tail()
```

Out[5]:

	Time	V1	V2	V3	V4	V5	V6	V7	
284802	172786.0	-11.881118	10.071785	-9.834783	-2.066656	-5.364473	-2.606837	-4.918215	7.30
284803	172787.0	-0.732789	-0.055080	2.035030	-0.738589	0.868229	1.058415	0.024330	0.29
284804	172788.0	1.919565	-0.301254	-3.249640	-0.557828	2.630515	3.031260	-0.296827	0.70
284805	172788.0	-0.240440	0.530483	0.702510	0.689799	-0.377961	0.623708	-0.686180	0.67
284806	172792.0	-0.533413	-0.189733	0.703337	-0.506271	-0.012546	-0.649617	1.577006	-0.41

5 rows × 31 columns



In the following cell we can see the datatypes of the features

In [6]:

```
data.info()
```

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 284807 entries, 0 to 284806
Data columns (total 31 columns):
#   Column  Non-Null Count  Dtype
---  -
0   Time    284807 non-null  float64
1   V1       284807 non-null  float64
2   V2       284807 non-null  float64
3   V3       284807 non-null  float64
4   V4       284807 non-null  float64
5   V5       284807 non-null  float64
6   V6       284807 non-null  float64
7   V7       284807 non-null  float64
8   V8       284807 non-null  float64
9   V9       284807 non-null  float64
10  V10      284807 non-null  float64
11  V11      284807 non-null  float64
12  V12      284807 non-null  float64
13  V13      284807 non-null  float64
14  V14      284807 non-null  float64
15  V15      284807 non-null  float64
16  V16      284807 non-null  float64
17  V17      284807 non-null  float64
18  V18      284807 non-null  float64
19  V19      284807 non-null  float64
20  V20      284807 non-null  float64
21  V21      284807 non-null  float64
```

```

22 V22      284807 non-null float64
23 V23      284807 non-null float64
24 V24      284807 non-null float64
25 V25      284807 non-null float64
26 V26      284807 non-null float64
27 V27      284807 non-null float64
28 V28      284807 non-null float64
29 Amount   284807 non-null float64
30 Class    284807 non-null int64
dtypes: float64(30), int64(1)
memory usage: 67.4 MB

```

Now we check for null values

```
In [7]: data.isnull().sum()
```

```

Out[7]: Time      0
V1         0
V2         0
V3         0
V4         0
V5         0
V6         0
V7         0
V8         0
V9         0
V10        0
V11        0
V12        0
V13        0
V14        0
V15        0
V16        0
V17        0
V18        0
V19        0
V20        0
V21        0
V22        0
V23        0
V24        0
V25        0
V26        0
V27        0
V28        0
Amount     0
Class      0
dtype: int64

```

We see that in this dataset, there are no null values anywhere.

```
In [8]: data=data.sample(frac=0.1, random_state=1)

print(data.shape)
```

```
(28481, 31)
```

We will be using about 10% of the dataset

Classifying the transactions into safe and fraudulent ones

Now we will check the column "Class" to see how many safe and fraudulent transactions there are in the dataset

```
In [9]: data["Class"].value_counts()
```

```
Out[9]: 0    28432
        1      49
        Name: Class, dtype: int64
```

```
In [10]: safe=data[data.Class==0]
        fraud=data[data.Class==1]
        print('No. of safe cases= {}'.format(len(safe)))
        print('No. of fraud cases= {}'.format(len(fraud)))

        frac=len(fraud)/len(safe)
        print('Fraction of fraud cases= {}'.format(frac))
```

```
No. of safe cases= 28432
No. of fraud cases= 49
Fraction of fraud cases= 0.0017234102419808666
```

We can see that the number of fraud transactions is very less compared to the number of safe transactions (only about 0.17%)

```
In [11]: print(safe.shape)
```

```
(28432, 31)
```

We have printed the shape of the safe transactions part of the dataset

Now we will see the shape of the fraud transactions part of the dataset

```
In [12]: print(fraud.shape)
```

```
(49, 31)
```

```
In [13]: safe.Amount.describe()
```

```
Out[13]: count    28432.000000
        mean      89.813898
        std       270.636594
        min        0.000000
        25%        5.990000
        50%       22.380000
        75%       78.820000
        max     19656.530000
        Name: Amount, dtype: float64
```

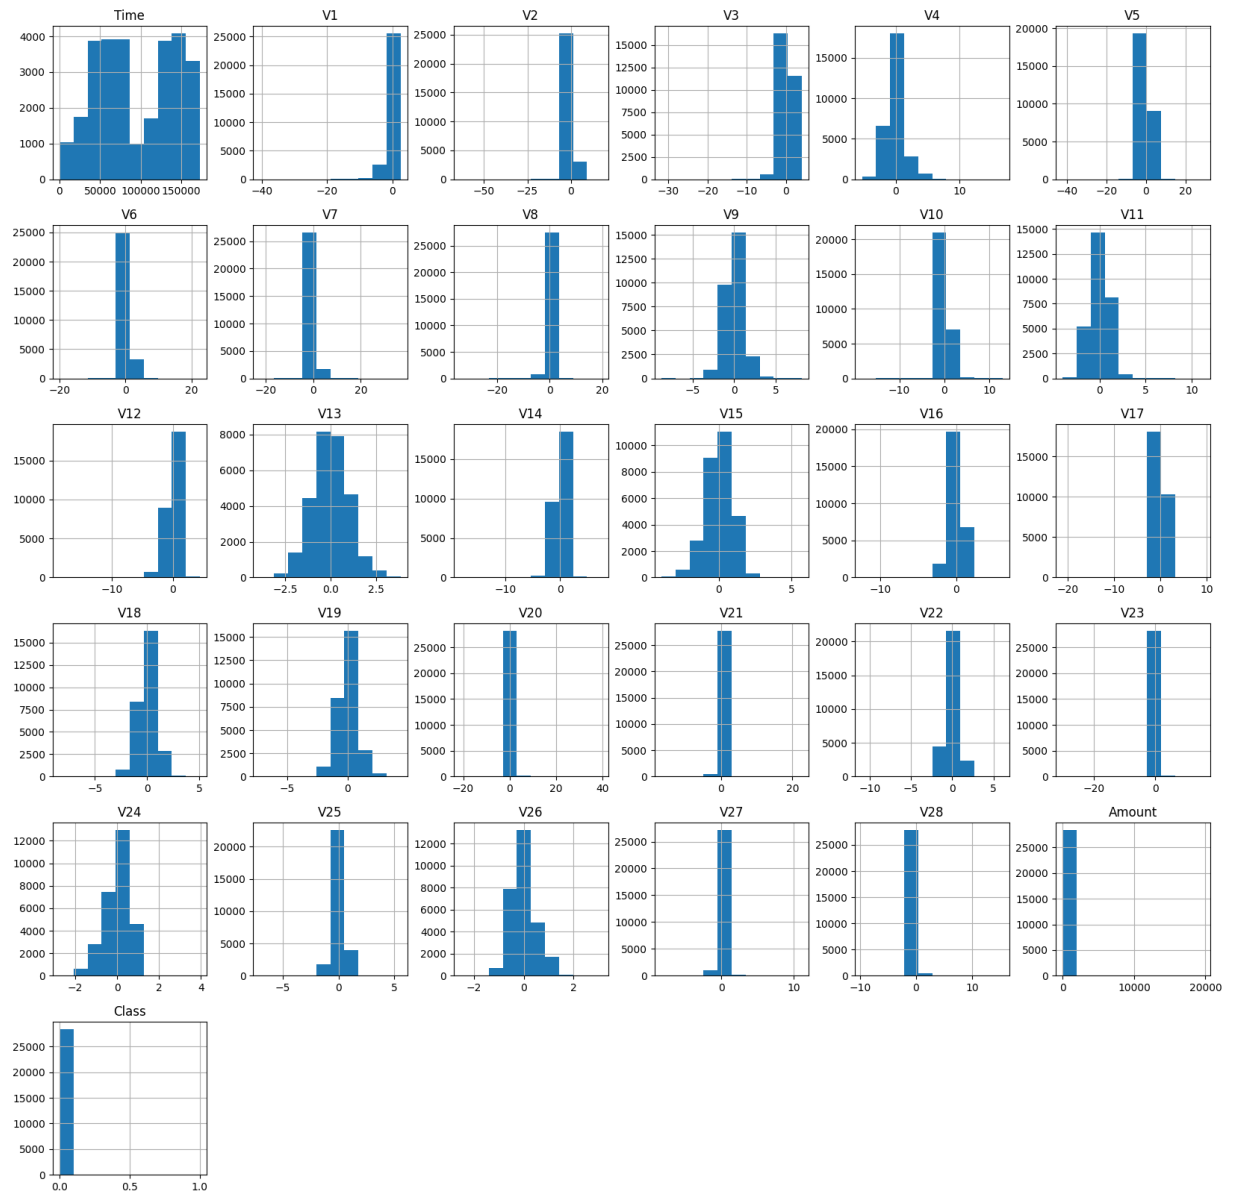
```
In [14]: fraud.Amount.describe()
```

```
Out[14]: count      49.000000
        mean     173.505306
        std     387.996569
        min      0.000000
        25%      1.000000
        50%      4.900000
        75%    122.680000
        max    2125.870000
        Name: Amount, dtype: float64
```

Visualising the dataset

```
In [15]:
```


```
data.hist(figsize=(20,20))
plt.show()
```



Here we have printed the histograms for the dataset we are using. We can see in the histogram for the "Class" column that the number of fraudulent transactions is negligible compared to the safe transactions


Next we look at a relational plot

```
In [16]: sns.relplot(x='Amount', y='Time', hue='Class', data=data)
plt.show()
```



Now, we make a correlation matrix, and then a heatmap based on that

```
In [17]: cormat=data.corr()  
fig=plt.figure(figsize=(12,8))  
  
sns.heatmap(cormat, square=True)  
plt.show()
```



Importing the libraries for processing our dataset

```
In [18]: from sklearn import linear_model
from sklearn.model_selection import train_test_split
```

Segregating the data into features and target

```
In [19]: X=data.iloc[:, :-1]
y=data["Class"]
```

```
In [20]: X_train, X_test, y_train, y_test=train_test_split(X, y, test_size=0.2)
```

We have taken our test size as 20% of the dataset

```
In [21]: print(X.shape)
print(y.shape)
```

```
(28481, 30)
(28481,)
```

Here we can see that the dataset has been split into the features and target

Importing our classifiers

We will now import the IsolationForest classifier, and the LocalOutlierFactor classifier. They are used for anomaly detection.

The LocalOutlierFactor calculates the anomaly score of each sample. It measures the local deviation of density of a given sample with respect to its neighbours.

The IsolationForest isolates the observations by randomly selecting a feature and then randomly selecting a value between the minimum and maximum values of that select feature. It returns an anomaly score.

```
In [22]: from sklearn.metrics import classification_report, accuracy_score
        from sklearn.ensemble import IsolationForest, RandomForestClassifier
        from sklearn.neighbors import LocalOutlierFactor
```

Defining a random state

```
In [23]: state=1
```

Defining the outlier detection methods

```
In [24]: classifiers={
        "Isolation Forest": IsolationForest(max_samples=len(X),
                                             contamination=frac,
                                             random_state=state),
        "Local Outlier Factor": LocalOutlierFactor(n_neighbors=20,
                                                  contamination=frac),
        "Random Forest": RandomForestClassifier(max_samples=len(X),
                                                random_state=state)
    }
```

Fitting the model

```
In [25]: n_outliers=len(fraud)

        for i, (clf_name, clf) in enumerate(classifiers.items()):
            #fit data and tag outliers
            if clf_name=='Local Outlier Factor':
                y_pred=clf.fit_predict(X)
                score_pred=clf.negative_outlier_factor_
            elif clf_name=='Isolation Forest':
                clf.fit(X)
                scores_pred=clf.decision_function(X)
                y_pred=clf.fit_predict(X)
            else:
                clf.fit(X, y)
                y_pred=clf.predict(X)

            #reshape pred vals to 0 for valid, 1 for fraud
            y_pred[y_pred==1] = 0
            y_pred[y_pred==-1]= 1

            n_errors=(y_pred != y).sum()

            #run classification metrics
            print("{}:{}".format(clf_name, n_errors))
            print(accuracy_score(y, y_pred))
            print(classification_report(y, y_pred))
```

Isolation Forest:71

0.99750711000316

	precision	recall	f1-score	support
0	1.00	1.00	1.00	28432
1	0.28	0.29	0.28	49
accuracy			1.00	28481
macro avg	0.64	0.64	0.64	28481
weighted avg	1.00	1.00	1.00	28481

Local Outlier Factor:97

0.9965942207085425

	precision	recall	f1-score	support
0	1.00	1.00	1.00	28432
1	0.02	0.02	0.02	49
accuracy			1.00	28481
macro avg	0.51	0.51	0.51	28481
weighted avg	1.00	1.00	1.00	28481

Random Forest:49

0.9982795547909132

	precision	recall	f1-score	support
0	1.00	1.00	1.00	28432
1	0.00	0.00	0.00	49
accuracy			1.00	28481
macro avg	0.50	0.50	0.50	28481
weighted avg	1.00	1.00	1.00	28481
